

```
In [1]: import pandas as pd
import os, shutil
import time
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam, Nadam, SGD
from tensorflow.keras import regularizers
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Ba
from tensorflow.keras.layers import Conv2D, MaxPooling2D, AveragePooling2D
from tensorflow.keras import losses

print(tf.__version__)    #YOUR TF VERSION CHECK
```

2.9.0

```
In [ ]: #Train data has folders for each class but Validation has just one folder.
#This section of the code iterates over the Val_annots file, creates folder

df = pd.read_csv('tiny-imagenet-200/val/Val_annots.csv')

parent_dir = 'tiny-imagenet-200/val/'
for i in range(len(df)):
    image_path = os.path.join(parent_dir, df.loc[i]['Name'])
    if not os.path.exists(parent_dir + df.loc[i]['Name']):
        os.mkdir(image_path)
    shutil.copy(parent_dir + 'images/' + df.loc[i]['Image'], image_path)

shutil.rmtree(parent_dir + 'images')
```

The Tiny ImageNet Dataset can be downloaded from <https://image-net.org/download-images.php> (<https://image-net.org/download-images.php>)

```
In [2]: # Define the image size and number of classes
img_size = 64
num_classes = 200

# Create the data generators for training and validation sets
train_datagen = ImageDataGenerator(
    rescale=1./255,
    horizontal_flip=True,
)

train_generator = train_datagen.flow_from_directory(
    'tiny-imagenet-200/train',
    target_size=(img_size, img_size),
    batch_size=64,
    class_mode='categorical'
)

val_datagen = ImageDataGenerator(
    rescale=1./255,
    horizontal_flip=True,
)

val_generator = val_datagen.flow_from_directory(
    'tiny-imagenet-200/val',
    target_size=(img_size, img_size),
    batch_size=64,
    class_mode='categorical'
)
```


Found 100000 images belonging to 200 classes.

Found 10000 images belonging to 200 classes.

```
In [3]: import numpy as np
from PIL import Image
import matplotlib.pyplot as plt

#Displaying a sample train image of Fish category
image = Image.open('tiny-imagenet-200/train/n01443537/images/n01443537_4.JPG')
image_array = np.asarray(image)

plt.imshow(image_array)
plt.show()
```



As we can see the image is only of 64X64X3 size and is highly blurred. Classification can be performed but expecting high validation accuracy can be challenging.

```
In [4]: #Below class is to record the model training times
class TimeHistory(tf.keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        self.times = []

    def on_epoch_begin(self, batch, logs={}):
        self.epoch_time_start = time.time()

    def on_epoch_end(self, batch, logs={}):
        epoch_time_end = time.time()
        self.times.append(epoch_time_end - self.epoch_time_start)
```

DEEP CNN Models

1. Single Convolutional Layer NN model

```
In [5]: # Define the CNN architecture
model1 = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(img_
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(1024, activation='relu'),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```

```
In [6]: print(model1.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
flatten (Flatten)	(None, 30752)	0
dense (Dense)	(None, 1024)	31491072
dense_1 (Dense)	(None, 200)	205000

```
=====
Total params: 31,696,968
Trainable params: 31,696,968
Non-trainable params: 0
```

None

```
In [7]: time_callback = TimeHistory()

# Compile the model with an optimizer, a loss function, and evaluation metrics
model1.compile(optimizer=Adam(learning_rate = 0.0001),
               loss='categorical_crossentropy',
               metrics=['accuracy'])

# Train the model on the training data and validate it on the validation data
with tf.device('/device:GPU:0'):
    history = model1.fit(train_generator,
                        validation_data=val_generator,
                        epochs=15,
                        callbacks=[time_callback, tf.keras.callbacks.TensorBoard])
```

Epoch 1/15

6/1563 [.....] - ETA: 36s - loss: 5.3947 - accuracy: 0.0104 WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0093s vs `on_train_batch_end` time: 0.0169s). Check your callbacks.

1563/1563 [=====] - 38s 24ms/step - loss: 4.7895 - accuracy: 0.0600 - val_loss: 4.4842 - val_accuracy: 0.0993

Epoch 2/15

1563/1563 [=====] - 36s 23ms/step - loss: 4.3215 - accuracy: 0.1191 - val_loss: 4.2536 - val_accuracy: 0.1265

Epoch 3/15

1563/1563 [=====] - 36s 23ms/step - loss: 4.0804 - accuracy: 0.1529 - val_loss: 4.1041 - val_accuracy: 0.1457

Epoch 4/15

1563/1563 [=====] - 36s 23ms/step - loss: 3.8904 - accuracy: 0.1814 - val_loss: 3.9833 - val_accuracy: 0.1618

Epoch 5/15

1563/1563 [=====] - 36s 23ms/step - loss: 3.7341 - accuracy: 0.2067 - val_loss: 3.8976 - val_accuracy: 0.1720

Epoch 6/15

1563/1563 [=====] - 36s 23ms/step - loss: 3.5990 - accuracy: 0.2288 - val_loss: 3.8270 - val_accuracy: 0.1868

Epoch 7/15

1563/1563 [=====] - 36s 23ms/step - loss: 3.4708 - accuracy: 0.2517 - val_loss: 3.7874 - val_accuracy: 0.1942

Epoch 8/15

1563/1563 [=====] - 37s 24ms/step - loss: 3.3522 - accuracy: 0.2710 - val_loss: 3.7486 - val_accuracy: 0.1979

Epoch 9/15

1563/1563 [=====] - 37s 24ms/step - loss: 3.2396 - accuracy: 0.2926 - val_loss: 3.7062 - val_accuracy: 0.2034

Epoch 10/15

1563/1563 [=====] - 37s 24ms/step - loss: 3.1341 - accuracy: 0.3120 - val_loss: 3.7092 - val_accuracy: 0.2050

Epoch 11/15

1563/1563 [=====] - 38s 24ms/step - loss: 3.0311 - accuracy: 0.3301 - val_loss: 3.6921 - val_accuracy: 0.2084

Epoch 12/15

1563/1563 [=====] - 37s 24ms/step - loss: 2.9279 - accuracy: 0.3504 - val_loss: 3.6742 - val_accuracy: 0.2143

Epoch 13/15

1563/1563 [=====] - 37s 24ms/step - loss: 2.8303 - accuracy: 0.3698 - val_loss: 3.6729 - val_accuracy: 0.2114

Epoch 14/15

1563/1563 [=====] - 37s 23ms/step - loss: 2.7354 - accuracy: 0.3883 - val_loss: 3.6693 - val_accuracy: 0.2207

Epoch 15/15

1563/1563 [=====] - 36s 23ms/step - loss: 2.6385 - accuracy: 0.4084 - val_loss: 3.6795 - val_accuracy: 0.2143

```
In [8]: total_time = sum(time_callback.times)
print(f"Total execution time: {total_time:.2f} seconds")
```

Total execution time: 552.10 seconds

2. Two Convolutional Layer NN model

```
In [9]: # Define the CNN architecture
model2 = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(img_
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(1024, activation='relu'),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```

```
In [10]: print(model2.summary())
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d_1 (MaxPooling 2D)	(None, 31, 31, 32)	0
conv2d_2 (Conv2D)	(None, 29, 29, 64)	18496
max_pooling2d_2 (MaxPooling 2D)	(None, 14, 14, 64)	0
flatten_1 (Flatten)	(None, 12544)	0
dense_2 (Dense)	(None, 1024)	12846080
dense_3 (Dense)	(None, 200)	205000
=====		
Total params: 13,070,472		
Trainable params: 13,070,472		
Non-trainable params: 0		
None		

```
In [11]: time_callback = TimeHistory()

# Compile the model with an optimizer, a loss function, and evaluation metrics
model2.compile(optimizer=Adam(learning_rate=0.0001),
               loss='categorical_crossentropy',
               metrics=['accuracy'])

# Train the model on the training data and validate it on the validation data
with tf.device('/device:GPU:0'):
    history = model2.fit(train_generator,
                        validation_data=val_generator,
                        epochs=15,
                        callbacks=[time_callback, tf.keras.callbacks.TensorBoard])
```



```

Epoch 1/15
  6/1563 [.....] - ETA: 41s - loss: 5.7219 - ac
curacy: 0.0026 WARNING:tensorflow:Callback method `on_train_batch_end`
is slow compared to the batch time (batch time: 0.0118s vs `on_train_batc
h_end` time: 0.0183s). Check your callbacks.
1563/1563 [=====] - 31s 20ms/step - loss: 4.9267
- accuracy: 0.0650 - val_loss: 4.5511 - val_accuracy: 0.1037
Epoch 2/15
1563/1563 [=====] - 34s 22ms/step - loss: 4.3696
- accuracy: 0.1274 - val_loss: 4.2424 - val_accuracy: 0.1458
Epoch 3/15
1563/1563 [=====] - 32s 21ms/step - loss: 4.0945
- accuracy: 0.1645 - val_loss: 4.0543 - val_accuracy: 0.1671
Epoch 4/15
1563/1563 [=====] - 31s 20ms/step - loss: 3.8943
- accuracy: 0.1944 - val_loss: 3.9258 - val_accuracy: 0.1874
Epoch 5/15
1563/1563 [=====] - 31s 20ms/step - loss: 3.7441
- accuracy: 0.2176 - val_loss: 3.8419 - val_accuracy: 0.1994
Epoch 6/15
1563/1563 [=====] - 33s 21ms/step - loss: 3.6196
- accuracy: 0.2376 - val_loss: 3.7395 - val_accuracy: 0.2154
Epoch 7/15
1563/1563 [=====] - 33s 21ms/step - loss: 3.5122
- accuracy: 0.2550 - val_loss: 3.6906 - val_accuracy: 0.2245
Epoch 8/15
1563/1563 [=====] - 32s 20ms/step - loss: 3.4221
- accuracy: 0.2710 - val_loss: 3.6632 - val_accuracy: 0.2292
Epoch 9/15
1563/1563 [=====] - 32s 20ms/step - loss: 3.3374
- accuracy: 0.2852 - val_loss: 3.6227 - val_accuracy: 0.2397
Epoch 10/15
1563/1563 [=====] - 32s 20ms/step - loss: 3.2593
- accuracy: 0.2984 - val_loss: 3.6215 - val_accuracy: 0.2384
Epoch 11/15
1563/1563 [=====] - 32s 20ms/step - loss: 3.1897
- accuracy: 0.3123 - val_loss: 3.5856 - val_accuracy: 0.2469
Epoch 12/15
1563/1563 [=====] - 32s 20ms/step - loss: 3.1169
- accuracy: 0.3241 - val_loss: 3.5701 - val_accuracy: 0.2502
Epoch 13/15
1563/1563 [=====] - 32s 20ms/step - loss: 3.0455
- accuracy: 0.3387 - val_loss: 3.5597 - val_accuracy: 0.2494
Epoch 14/15
1563/1563 [=====] - 32s 20ms/step - loss: 2.9837
- accuracy: 0.3494 - val_loss: 3.5479 - val_accuracy: 0.2540
Epoch 15/15
1563/1563 [=====] - 32s 20ms/step - loss: 2.9102
- accuracy: 0.3644 - val_loss: 3.5542 - val_accuracy: 0.2500

```

```

In [12]: total_time = sum(time_callback.times)
print(f"Total execution time: {total_time:.2f} seconds")

```

Total execution time: 478.81 seconds

3. DEEP CNN Model 1

```

In [13]: model3 = Sequential()

model3.add(Conv2D(64, (3, 3), strides=(1,1), padding='same', input_shape=(i
model3.add(BatchNormalization())
model3.add(Conv2D(64, (3, 3), strides=(1,1), padding='same', kernel_regular
model3.add(BatchNormalization())
model3.add(MaxPooling2D(pool_size=(3, 3), strides=(2, 2), padding='same'))
model3.add(Activation('relu'))

model3.add(Conv2D(128, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Conv2D(128, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(MaxPooling2D(pool_size=(3, 3), strides=(2, 2), padding='same'))
model3.add(Activation('relu'))

model3.add(Conv2D(128, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Conv2D(128, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Activation('relu'))
model3.add(MaxPooling2D(pool_size=(3, 3), strides=(2, 2), padding='same'))

model3.add(Conv2D(128, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Conv2D(128, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Activation('relu'))
model3.add(MaxPooling2D(pool_size=(3, 3), strides=(2, 2), padding='same'))

model3.add(Conv2D(256, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Conv2D(512, (3, 3), strides=(1,1), padding='same', kernel_regula
model3.add(BatchNormalization())
model3.add(Activation('relu'))
model3.add(MaxPooling2D(pool_size=(3, 3), strides=(2, 2), padding='same'))

model3.add(Flatten())
model3.add(Dense(4096))
model3.add(BatchNormalization())
model3.add(Activation('relu'))

model3.add(Dense(1024))
model3.add(BatchNormalization())
model3.add(Activation('relu'))

model3.add(Dense(num_classes))

model3.add(Activation('softmax'))

time_callback = TimeHistory()

model3.compile(loss=losses.categorical_crossentropy,
optimizer=Adam(learning_rate=0.0001), metrics=['accuracy'])

```

```
In [14]: print(model3.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 64, 64, 64)	1792
batch_normalization (Batch Normalization)	(None, 64, 64, 64)	256
conv2d_4 (Conv2D)	(None, 64, 64, 64)	36928
batch_normalization_1 (Batch Normalization)	(None, 64, 64, 64)	256
max_pooling2d_3 (MaxPooling2D)	(None, 32, 32, 64)	0
activation (Activation)	(None, 32, 32, 64)	0
conv2d_5 (Conv2D)	(None, 32, 32, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 32, 32, 128)	512
conv2d_6 (Conv2D)	(None, 32, 32, 128)	147584
batch_normalization_3 (Batch Normalization)	(None, 32, 32, 128)	512
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 128)	0
activation_1 (Activation)	(None, 16, 16, 128)	0
conv2d_7 (Conv2D)	(None, 16, 16, 128)	147584
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 128)	512
conv2d_8 (Conv2D)	(None, 16, 16, 128)	147584
batch_normalization_5 (Batch Normalization)	(None, 16, 16, 128)	512
activation_2 (Activation)	(None, 16, 16, 128)	0
max_pooling2d_5 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_9 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_6 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_10 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_7 (Batch Normalization)	(None, 8, 8, 128)	512

```

hNormalization)

activation_3 (Activation)      (None, 8, 8, 128)          0

max_pooling2d_6 (MaxPooling   (None, 4, 4, 128)          0
2D)

conv2d_11 (Conv2D)            (None, 4, 4, 256)          295168

batch_normalization_8 (Batc   (None, 4, 4, 256)          1024
hNormalization)

conv2d_12 (Conv2D)            (None, 4, 4, 512)          1180160

batch_normalization_9 (Batc   (None, 4, 4, 512)          2048
hNormalization)

activation_4 (Activation)      (None, 4, 4, 512)          0

max_pooling2d_7 (MaxPooling   (None, 2, 2, 512)          0
2D)

flatten_2 (Flatten)           (None, 2048)                0

dense_4 (Dense)                (None, 4096)                8392704

batch_normalization_10 (Bat   (None, 4096)                16384
chNormalization)

activation_5 (Activation)      (None, 4096)                0

dense_5 (Dense)                (None, 1024)               4195328

batch_normalization_11 (Bat   (None, 1024)               4096
chNormalization)

activation_6 (Activation)      (None, 1024)                0

dense_6 (Dense)                (None, 200)                205000

activation_7 (Activation)      (None, 200)                0

=====
Total params: 15,145,992
Trainable params: 15,132,424
Non-trainable params: 13,568

None

```

```
In [15]: # Train the model on the training data and validate it on the validation data
with tf.device('/device:GPU:0'):
    history = model3.fit(train_generator,
                        validation_data=val_generator,
                        epochs=15,
                        callbacks=[time_callback, tf.keras.callbacks.TensorBoard])
```

Epoch 1/15

6/1563 [.....] - ETA: 3:42 - loss: 9.4461 - accuracy: 0.0052WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0620s vs `on_train_batch_end` time: 0.1052s). Check your callbacks.

1563/1563 [=====] - 197s 125ms/step - loss: 7.1246 - accuracy: 0.1654 - val_loss: 6.1606 - val_accuracy: 0.2170

Epoch 2/15

1563/1563 [=====] - 196s 125ms/step - loss: 5.3315 - accuracy: 0.2849 - val_loss: 4.9999 - val_accuracy: 0.2842

Epoch 3/15

1563/1563 [=====] - 196s 125ms/step - loss: 4.3545 - accuracy: 0.3569 - val_loss: 4.3981 - val_accuracy: 0.3104

Epoch 4/15

1563/1563 [=====] - 196s 125ms/step - loss: 3.7288 - accuracy: 0.4108 - val_loss: 4.1502 - val_accuracy: 0.3201

Epoch 5/15

1563/1563 [=====] - 196s 125ms/step - loss: 3.2835 - accuracy: 0.4561 - val_loss: 3.6812 - val_accuracy: 0.3706

Epoch 6/15

1563/1563 [=====] - 196s 125ms/step - loss: 2.9388 - accuracy: 0.4983 - val_loss: 3.4901 - val_accuracy: 0.3859

Epoch 7/15

1563/1563 [=====] - 196s 125ms/step - loss: 2.6481 - accuracy: 0.5429 - val_loss: 3.3602 - val_accuracy: 0.4059

Epoch 8/15

1563/1563 [=====] - 208s 133ms/step - loss: 2.3932 - accuracy: 0.5840 - val_loss: 3.3655 - val_accuracy: 0.3955

Epoch 9/15

1563/1563 [=====] - 210s 135ms/step - loss: 2.1591 - accuracy: 0.6242 - val_loss: 3.3649 - val_accuracy: 0.3959

Epoch 10/15

1563/1563 [=====] - 210s 135ms/step - loss: 1.9417 - accuracy: 0.6687 - val_loss: 3.3642 - val_accuracy: 0.3958

Epoch 11/15

1563/1563 [=====] - 210s 135ms/step - loss: 1.7443 - accuracy: 0.7098 - val_loss: 3.3743 - val_accuracy: 0.3960

Epoch 12/15

1563/1563 [=====] - 211s 135ms/step - loss: 1.5586 - accuracy: 0.7505 - val_loss: 3.3587 - val_accuracy: 0.4036

Epoch 13/15

1563/1563 [=====] - 211s 135ms/step - loss: 1.4000 - accuracy: 0.7860 - val_loss: 3.5995 - val_accuracy: 0.3777

Epoch 14/15

1563/1563 [=====] - 211s 135ms/step - loss: 1.2616 - accuracy: 0.8176 - val_loss: 3.4843 - val_accuracy: 0.3940

Epoch 15/15

1563/1563 [=====] - 211s 135ms/step - loss: 1.1472 - accuracy: 0.8426 - val_loss: 3.5772 - val_accuracy: 0.3954

```
In [16]: total_time = sum(time_callback.times)
print(f"Total execution time: {total_time:.2f} seconds")
```

Total execution time: 3054.00 seconds

4. Deep CNN Model 2

```
In [17]: # Import required libraries
import tensorflow as tf
from tensorflow.keras import layers

# Define the input shape
input_shape = (64, 64, 3)

# Define the model layers
model4 = tf.keras.Sequential([
    layers.Conv2D(64, (3,3), padding='same', activation='relu', input_shape
    layers.Conv2D(64, (3,3), padding='same', activation='relu', kernel_regu
    layers.MaxPooling2D((2,2)),
    layers.BatchNormalization(),

    layers.Conv2D(128, (3,3), padding='same', activation='relu', kernel_reg
    layers.MaxPooling2D((2,2)),
    layers.BatchNormalization(),

    layers.Conv2D(256, (3,3), padding='same', activation='relu', kernel_reg
    layers.Conv2D(256, (3,3), padding='same', activation='relu', kernel_reg
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),

    layers.Dense(1024, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(512, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(200, activation='softmax')
])

time_callback = TimeHistory()

# Compile the model
model4.compile(optimizer=Adam(learning_rate=0.0001),
               loss='categorical_crossentropy',
               metrics=['accuracy'])
```



```
In [18]: # Print the model summary
print(model4.summary())
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_13 (Conv2D)	(None, 64, 64, 64)	1792
conv2d_14 (Conv2D)	(None, 64, 64, 64)	36928
max_pooling2d_8 (MaxPooling2D)	(None, 32, 32, 64)	0
batch_normalization_12 (Batch Normalization)	(None, 32, 32, 64)	256
conv2d_15 (Conv2D)	(None, 32, 32, 128)	73856
max_pooling2d_9 (MaxPooling2D)	(None, 16, 16, 128)	0
batch_normalization_13 (Batch Normalization)	(None, 16, 16, 128)	512
conv2d_16 (Conv2D)	(None, 16, 16, 256)	295168
conv2d_17 (Conv2D)	(None, 16, 16, 256)	590080
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 256)	0
flatten_3 (Flatten)	(None, 16384)	0
dense_7 (Dense)	(None, 1024)	16778240
dropout (Dropout)	(None, 1024)	0
dense_8 (Dense)	(None, 512)	524800
dropout_1 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 200)	102600
=====		
Total params: 18,404,232		
Trainable params: 18,403,848		
Non-trainable params: 384		

None

```
In [19]: # Train the model
with tf.device('/device:GPU:0'):
    history = model4.fit(train_generator,
                        epochs=25,
                        validation_data=val_generator,
                        callbacks=[time_callback, tf.keras.callbacks.Tensor
```

Epoch 1/25

6/1563 [.....] - ETA: 2:26 - loss: 11.5748 - accuracy: 0.0052 WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0323s vs `on_train_batch_end` time: 0.0623s). Check your callbacks.

1563/1563 [=====] - 151s 96ms/step - loss: 7.1938 - accuracy: 0.0339 - val_loss: 5.0122 - val_accuracy: 0.1006

Epoch 2/25

1563/1563 [=====] - 150s 96ms/step - loss: 4.6325 - accuracy: 0.1095 - val_loss: 4.0970 - val_accuracy: 0.1766

Epoch 3/25

1563/1563 [=====] - 149s 96ms/step - loss: 4.0229 - accuracy: 0.1744 - val_loss: 3.6336 - val_accuracy: 0.2409

Epoch 4/25

1563/1563 [=====] - 149s 96ms/step - loss: 3.6961 - accuracy: 0.2189 - val_loss: 3.4759 - val_accuracy: 0.2642

Epoch 5/25

1563/1563 [=====] - 149s 96ms/step - loss: 3.4683 - accuracy: 0.2555 - val_loss: 3.2908 - val_accuracy: 0.2893

Epoch 6/25

1563/1563 [=====] - 150s 96ms/step - loss: 3.2929 - accuracy: 0.2885 - val_loss: 3.1674 - val_accuracy: 0.3149

Epoch 7/25

1563/1563 [=====] - 150s 96ms/step - loss: 3.1501 - accuracy: 0.3148 - val_loss: 3.0611 - val_accuracy: 0.3317

Epoch 8/25

1563/1563 [=====] - 150s 96ms/step - loss: 3.0179 - accuracy: 0.3410 - val_loss: 2.9928 - val_accuracy: 0.3503

Epoch 9/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.9114 - accuracy: 0.3603 - val_loss: 2.9234 - val_accuracy: 0.3588

Epoch 10/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.8053 - accuracy: 0.3821 - val_loss: 2.8927 - val_accuracy: 0.3678

Epoch 11/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.7158 - accuracy: 0.3999 - val_loss: 2.8264 - val_accuracy: 0.3791

Epoch 12/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.6325 - accuracy: 0.4184 - val_loss: 2.7976 - val_accuracy: 0.3892

Epoch 13/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.5517 - accuracy: 0.4352 - val_loss: 2.7569 - val_accuracy: 0.4004

Epoch 14/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.4762 - accuracy: 0.4508 - val_loss: 2.8391 - val_accuracy: 0.3837

Epoch 15/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.3996 - accuracy: 0.4689 - val_loss: 2.7409 - val_accuracy: 0.3989

Epoch 16/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.3301 - accuracy: 0.4836 - val_loss: 2.7665 - val_accuracy: 0.4049

Epoch 17/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.2676 - accuracy: 0.4976 - val_loss: 2.7119 - val_accuracy: 0.4177

Epoch 18/25

1563/1563 [=====] - 150s 96ms/step - loss: 2.197

```

2 - accuracy: 0.5127 - val_loss: 2.6823 - val_accuracy: 0.4169
Epoch 19/25
1563/1563 [=====] - 150s 96ms/step - loss: 2.142
2 - accuracy: 0.5253 - val_loss: 2.7598 - val_accuracy: 0.4093
Epoch 20/25
1563/1563 [=====] - 150s 96ms/step - loss: 2.067
9 - accuracy: 0.5410 - val_loss: 2.7221 - val_accuracy: 0.4182
Epoch 21/25
1563/1563 [=====] - 150s 96ms/step - loss: 2.009
6 - accuracy: 0.5550 - val_loss: 2.8340 - val_accuracy: 0.3975
Epoch 22/25
1563/1563 [=====] - 150s 96ms/step - loss: 1.953
9 - accuracy: 0.5672 - val_loss: 2.8361 - val_accuracy: 0.4000
Epoch 23/25
1563/1563 [=====] - 149s 95ms/step - loss: 1.894
8 - accuracy: 0.5814 - val_loss: 2.7627 - val_accuracy: 0.4200
Epoch 24/25
1563/1563 [=====] - 149s 95ms/step - loss: 1.842
2 - accuracy: 0.5928 - val_loss: 2.7544 - val_accuracy: 0.4185
Epoch 25/25
1563/1563 [=====] - 150s 96ms/step - loss: 1.778
9 - accuracy: 0.6089 - val_loss: 2.7473 - val_accuracy: 0.4240

```

```

In [20]: total_time = sum(time_callback.times)
print(f"Total execution time: {total_time:.2f} seconds")

```

Total execution time: 3744.04 seconds

5) ResNet Model 1- ResNet50

```

In [5]: resnet50_model = tf.keras.applications.ResNet50(
        weights='imagenet',
        include_top=False,
        input_shape=(img_size, img_size, 3)
    )

```

```

In [6]: x = resnet50_model.output
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.5)(x)
predictions = layers.Dense(num_classes, activation='softmax')(x)

model5 = models.Model(inputs = resnet50_model.input, outputs=predictions)

```

```

In [7]: time_callback = TimeHistory()

# Compile the model
model5.compile(optimizer=Adam(learning_rate=0.0001),
               loss='categorical_crossentropy',
               metrics=['accuracy'])

```

```
In [8]: print(model5.summary())
```

```
conv2_block3_out (Activation) (None, 16, 16, 256) 0 ['conv2_
block3_add[0][0]']

conv3_block1_1_conv (Conv2D) (None, 8, 8, 128) 32896 ['conv2_
block3_out[0][0]']

conv3_block1_1_bn (BatchNormal (None, 8, 8, 128) 512 ['conv3_
block1_1_conv[0][0]']
ization)

conv3_block1_1_relu (Activatio (None, 8, 8, 128) 0 ['conv3_
block1_1_bn[0][0]']
n)

conv3_block1_2_conv (Conv2D) (None, 8, 8, 128) 147584 ['conv3_
block1_1_relu[0][0]']

conv3_block1_2_bn (BatchNormal (None, 8, 8, 128) 512 ['conv3_
block1_2_conv[0][0]']
```

```
In [9]: # Train the model
with tf.device('/device:GPU:0'):
    history = model5.fit(train_generator,
                        epochs=15,
                        validation_data=val_generator,
                        callbacks=[time_callback, tf.keras.callbacks.TensorBoard])
```

Epoch 1/15

6/1563 [.....] - ETA: 3:31 - loss: 6.8956 - accuracy: 0.0000e+00WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0998s vs `on_train_batch_end` time: 0.1228s). Check your callbacks.

1563/1563 [=====] - 228s 143ms/step - loss: 3.4713 - accuracy: 0.2689 - val_loss: 2.3051 - val_accuracy: 0.4529

Epoch 2/15

1563/1563 [=====] - 222s 142ms/step - loss: 2.1666 - accuracy: 0.4841 - val_loss: 2.0099 - val_accuracy: 0.5088

Epoch 3/15

1563/1563 [=====] - 221s 141ms/step - loss: 1.7259 - accuracy: 0.5727 - val_loss: 1.9419 - val_accuracy: 0.5334

Epoch 4/15

1563/1563 [=====] - 219s 140ms/step - loss: 1.4161 - accuracy: 0.6404 - val_loss: 1.9479 - val_accuracy: 0.5355

Epoch 5/15

1563/1563 [=====] - 220s 140ms/step - loss: 1.1795 - accuracy: 0.6917 - val_loss: 1.8960 - val_accuracy: 0.5548

Epoch 6/15

1563/1563 [=====] - 220s 141ms/step - loss: 0.9778 - accuracy: 0.7399 - val_loss: 1.9514 - val_accuracy: 0.5486

Epoch 7/15

1563/1563 [=====] - 220s 141ms/step - loss: 0.8171 - accuracy: 0.7795 - val_loss: 2.0127 - val_accuracy: 0.5491

Epoch 8/15

1563/1563 [=====] - 221s 142ms/step - loss: 0.6878 - accuracy: 0.8100 - val_loss: 2.0410 - val_accuracy: 0.5512

Epoch 9/15

1563/1563 [=====] - 221s 142ms/step - loss: 0.5861 - accuracy: 0.8357 - val_loss: 2.1752 - val_accuracy: 0.5480

Epoch 10/15

1563/1563 [=====] - 222s 142ms/step - loss: 0.5017 - accuracy: 0.8576 - val_loss: 2.2624 - val_accuracy: 0.5412

Epoch 11/15

1563/1563 [=====] - 221s 141ms/step - loss: 0.4399 - accuracy: 0.8738 - val_loss: 2.2292 - val_accuracy: 0.5575

Epoch 12/15

1563/1563 [=====] - 221s 141ms/step - loss: 0.3821 - accuracy: 0.8886 - val_loss: 2.2521 - val_accuracy: 0.5526

Epoch 13/15

1563/1563 [=====] - 221s 141ms/step - loss: 0.3424 - accuracy: 0.8990 - val_loss: 2.3511 - val_accuracy: 0.5531

Epoch 14/15

1563/1563 [=====] - 222s 142ms/step - loss: 0.3040 - accuracy: 0.9095 - val_loss: 2.4762 - val_accuracy: 0.5334

Epoch 15/15

1563/1563 [=====] - 222s 142ms/step - loss: 0.2901 - accuracy: 0.9126 - val_loss: 2.4808 - val_accuracy: 0.5395

```
In [10]: total_time = sum(time_callback.times)
print(f"Total execution time: {total_time:.2f} seconds")
```

Total execution time: 3320.05 seconds

6) ResNet Model 2 - ResNet101

```
In [11]: resnet101_model = tf.keras.applications.ResNet101(
    include_top=False,
    weights="imagenet",
    input_shape=(img_size, img_size, 3),
)
```

```
In [12]: x = resnet101_model.output
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.5)(x)
predictions = layers.Dense(num_classes, activation='softmax')(x)

model6 = models.Model(inputs = resnet101_model.input, outputs=predictions)
```

```
In [13]: time_callback = TimeHistory()

# Compile the model
model6.compile(optimizer=Adam(learning_rate=0.0001),
               loss='categorical_crossentropy',
               metrics=['accuracy'])
```

```
In [14]: print(model6.summary())
```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_2 (InputLayer)	[(None, 64, 64, 3)]	0	[]
conv1_pad (ZeroPadding2D)	(None, 70, 70, 3)	0	['input_2[0][0]']
conv1_conv (Conv2D)	(None, 32, 32, 64)	9472	['conv1_pad[0][0]']
conv1_bn (BatchNormalization)	(None, 32, 32, 64)	256	['conv1_conv[0][0]']
conv1_relu (Activation)	(None, 32, 32, 64)	0	['conv1_bn[0][0]']

```
In [15]: # Train the model
with tf.device('/device:GPU:0'):
    history = model6.fit(train_generator,
                        epochs=15,
                        validation_data=val_generator,
                        callbacks=[time_callback, tf.keras.callbacks.TensorBoard])
```

```
Epoch 1/15
1563/1563 [=====] - 381s 237ms/step - loss: 3.43
22 - accuracy: 0.2747 - val_loss: 2.2370 - val_accuracy: 0.4663
Epoch 2/15
1563/1563 [=====] - 365s 233ms/step - loss: 2.08
05 - accuracy: 0.5015 - val_loss: 2.0029 - val_accuracy: 0.5235
Epoch 3/15
1563/1563 [=====] - 363s 232ms/step - loss: 1.62
27 - accuracy: 0.5986 - val_loss: 1.8773 - val_accuracy: 0.5395
Epoch 4/15
1563/1563 [=====] - 362s 232ms/step - loss: 1.30
10 - accuracy: 0.6671 - val_loss: 1.9223 - val_accuracy: 0.5489
Epoch 5/15
1563/1563 [=====] - 362s 232ms/step - loss: 1.05
09 - accuracy: 0.7246 - val_loss: 1.9876 - val_accuracy: 0.5406
Epoch 6/15
1563/1563 [=====] - 364s 233ms/step - loss: 0.84
78 - accuracy: 0.7735 - val_loss: 2.0723 - val_accuracy: 0.5532
Epoch 7/15
1563/1563 [=====] - 362s 232ms/step - loss: 0.68
97 - accuracy: 0.8113 - val_loss: 2.0589 - val_accuracy: 0.5498
Epoch 8/15
1563/1563 [=====] - 363s 233ms/step - loss: 0.56
99 - accuracy: 0.8409 - val_loss: 2.1778 - val_accuracy: 0.5487
Epoch 9/15
1563/1563 [=====] - 362s 232ms/step - loss: 0.48
21 - accuracy: 0.8636 - val_loss: 2.1881 - val_accuracy: 0.5507
Epoch 10/15
1563/1563 [=====] - 363s 233ms/step - loss: 0.42
13 - accuracy: 0.8800 - val_loss: 2.2289 - val_accuracy: 0.5505
Epoch 11/15
1563/1563 [=====] - 364s 233ms/step - loss: 0.36
73 - accuracy: 0.8935 - val_loss: 2.3453 - val_accuracy: 0.5432
Epoch 12/15
1563/1563 [=====] - 361s 231ms/step - loss: 0.31
72 - accuracy: 0.9058 - val_loss: 2.4139 - val_accuracy: 0.5423
Epoch 13/15
1563/1563 [=====] - 361s 231ms/step - loss: 0.29
44 - accuracy: 0.9131 - val_loss: 2.5502 - val_accuracy: 0.5351
Epoch 14/15
1563/1563 [=====] - 366s 234ms/step - loss: 0.26
22 - accuracy: 0.9216 - val_loss: 2.4279 - val_accuracy: 0.5477
Epoch 15/15
1563/1563 [=====] - 363s 232ms/step - loss: 0.23
90 - accuracy: 0.9280 - val_loss: 2.5798 - val_accuracy: 0.5337
```



```
In [16]: total_time = sum(time_callback.times)
print(f"Total execution time: {total_time:.2f} seconds")
```

Total execution time: 5463.40 seconds

```
In [ ]:
```